

# Crash Course on UNIX and Systems Tools

---

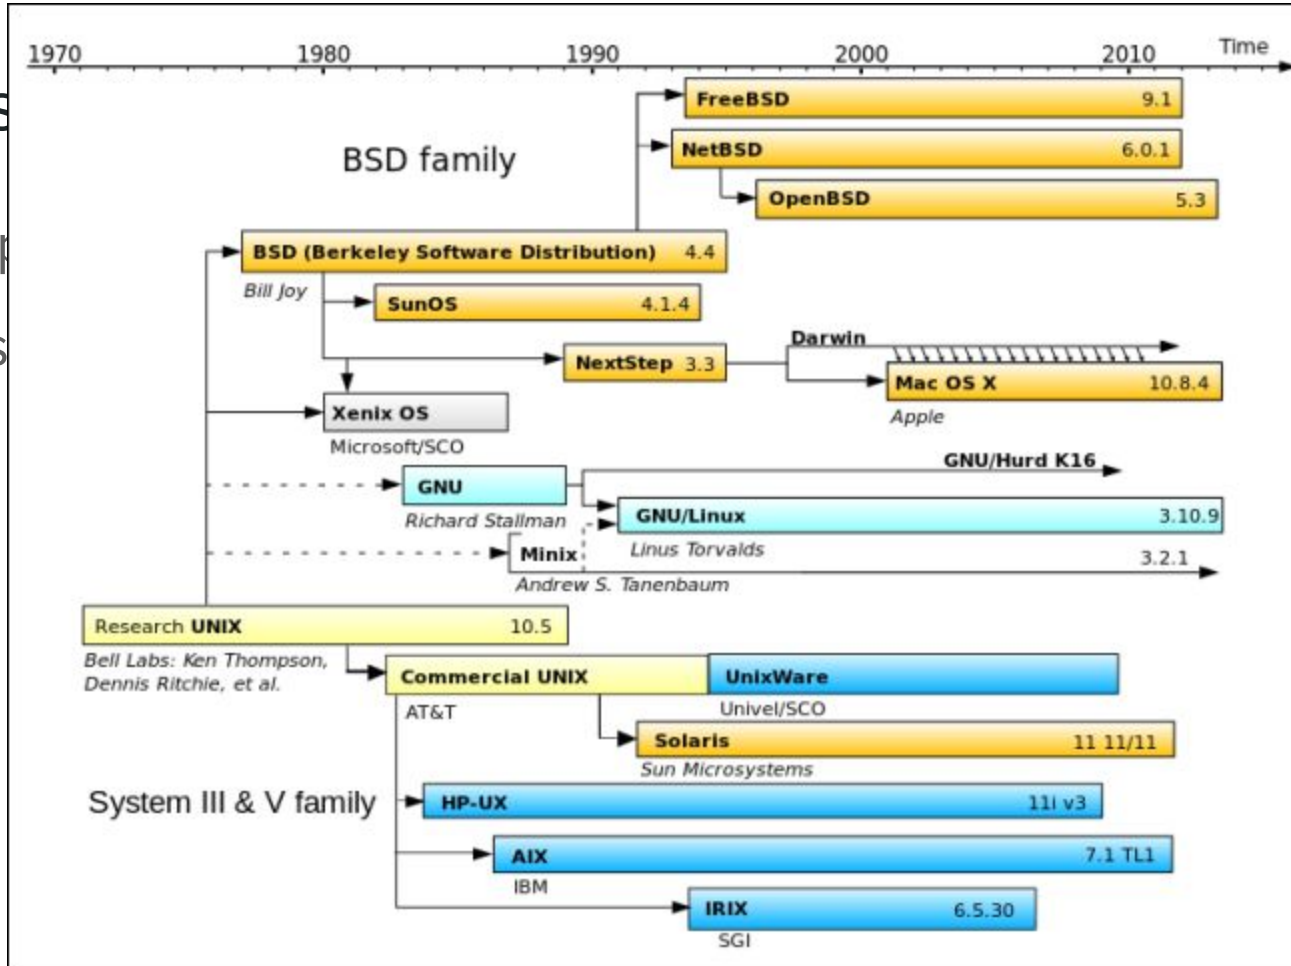
Day 1 --- The UNIX Environment and The Shell

# What is UNIX?

- An operating system from the 1960s
- Consists of a kernel, **shell**, and modules/programs to run
- Who cares? UNIX design is *very* influential

What is

- An op
- Cons
- Who



# What is UNIX?

- An operating system from the 1960s
- Consists of a kernel, **shell**, and modules/programs to run
- Who cares? UNIX design is *very* influential
- *Really* important base for many systems and low-level tools
- 213, 343, many other systems courses

# What is the shell?

- The interactive interface between the user and the OS --- you can run programs, run scripts, execute shell commands, write to files/directories, etc.
- Many brands of shells --- we're going to use **bash**

A black rectangular box representing a terminal window. Inside the box, on the left side, is a white dollar sign (\$) followed by a white vertical bar (|), which is a common prompt for the bash shell.

```
$ |
```

## Navigation: *Where am I?*

`pwd`

- “Present working directory”
- Outputs your current **absolute path**

-----

```
$ pwd
```

```
/root/sandbox
```

## Navigation: *Where am I?*

- **Absolute path:** location with respect to the **root directory** “/”
- **Relative path:** location with respect to:
  - The **home** directory “~/”
  - The **current** directory “./”
  - The **parent** directory “../”
  - Any combination of the above choices, i.e. “./mydir/../” etc.

## Navigation: *What's in here?*

**ls**

- “List”
- Outputs all files/directories inside the **current** directory
- Or, specify a path (absolute or relative) to see its contents

-----

```
$ ls
```

```
file1  file2
```



## Navigation: *What's in here?*

```
$ ls -a
```

```
./ ../ .bash_history another_file my_file
```

- Entries that start with “.”

## Navigation: *What's in here?*

```
$ ls -ltr
```

```
total 1
```

```
-rw-r--r-- 1 root root 0 Jan  2 23:07 file1
```

```
-rw-r--r-- 1 root root 0 Jan  2 23:07 file2
```

- **-ltr** --- Show details of entries in chronological order
  - **-l** lists in “**long**” format
  - **-t** sort entries by **time**
  - **-r** **reverse** the order

## Navigation: *What's in here?*

```
$ ls -ltr
```

```
total 1
```

```
-rw-r--r-- 1 root root 0 Jan  2 23:07 file1
```

```
-rw-r--r-- 1 root root 0 Jan  2 23:07 file2
```

- What does all of this info mean? **man-pages!**

```
$ man ls
```

- Spacebar to move down, “q” to quit

## Navigation: *Where can I go?*

**cd**

- “Change directory”
- Changes current directory based on specified path

-----

```
$ cd ../
```

```
Desktop/  Documents/  Downloads/  Music/
```

```
Pictures/  Public/  Templates/  Videos/  sandbox/
```

## Building: *Creating directories*

### **mkdir**

- “Make directory”
- Creates new directory based on specified name

-----

```
$ mkdir mydir
```

```
$ ls
```

```
file1  file2  mydir/
```

## Building: *Creating directories*

```
$ mkdir parent/child
```

```
mkdir: cannot create directory 'parent/child': No  
such file or directory
```

```
$ mkdir -p parent/child
```

- **-p** builds the specified directory structure if it doesn't already exist (including parent directories)
- Very useful when scripting

## Building: *Creating files*

### **touch**

- Modifies file timestamps --- creation, access, modification, etc.
- Often used to create files (without input or redirection)

-----

```
$ cd mydir ; touch newfile
```

```
$ ls
```

```
newfile
```

## Building: *Moving and renaming files*

**mv**

- “Move”
- Moves or renames files/directories (**be careful!** --- see manual)

-----

```
$ mv newfile otherfile
```

```
$ ls
```

```
otherfile
```



## Building: *Moving and renaming files*

```
$ cd ../ ; ls
file1  file2  mydir/
$ mv mydir/ otherdir
$ ls
file1  file2  otherdir/
$ mv file1 otherdir/ ; ls
file2  otherdir/
$ ls otherdir/
file1  otherfile
```

## Building: *Removing files*

**rm**

- “Remove”
- Removes files (**be careful!** --- see manual)

-----

```
$ rm file2 ; ls  
otherdir/
```

## Building: *Removing directories*

### `rmdir`

- “Remove directory”
- Removes the specified directory if it's empty

-----

```
$ rmdir otherdir/
```

```
rmdir: failed to remove 'otherdir/': Directory  
not empty
```

## Building: *Removing directories*

```
$ rm -r mydir
```

- Removes a directory *and* its files recursively

-----

```
$ ls otherdir/
```

```
file1  otherfile
```

```
$ rm -r otherdir/
```

```
$ ls
```

## Building: *Redirection*

“>” (redirect stdout)

“&>” (redirect stdout and stderr)

“>>” (append stdout to ...)

“&>>” (append stdout and stderr to ...)

-----

```
$ lscpu > lscpu.out
```

## Building: *Viewing files*

more and **less**

- Utilities to view or **page** through files
- **less** is a modern version of **more**, please use **less**
- <http://www.greenwoodsoftware.com/less/faq.html#history>

-----

**\$ less lscpu.out**

- “Space” to page, can use “u” , “d” , or arrows to move

## Building: *Writing to files*

### echo

- Output a string to the shell (**stdout**)
- **-e** recognizes escaped characters

-----

```
$ echo "Hello world!"
```

```
Hello world!
```

```
$ echo -e "\"Hello world\"!"
```

```
"Hello world"!
```

## Building: *Writing to files*

```
$ echo "Hello world!" > hello.out
```

```
$ less hello.out
```

```
Hello world!
```

- Echo a string and redirect it to a file
  - “>” operator b/c **stdout**
  - View the file (via **less**)



## Building: *Writing to files*

### `cat`

- Concatenate files, write to **`stdout`**, and even write to files
- If no path is given, **`cat`** begins to read from **`stdin`**

-----

```
$ cat hello.out  
Hello world!
```

## Building: *Writing to files*

```
$ cat > hello2.out
```

```
This is some text
```

```
^D
```

- Writing to a file requires redirection (“>”) to a file (**hello2.out**)
- Enter **ctrl d** (^D) to finish writing

## Building: *Writing to files*

```
$ cat hello.out hello2.out > hellos.out
```

```
$ cat hellos.out
```

```
Hello world!
```

```
This is some text
```

- Specify however many files --- redirect them to an output file

## Building: *Revisiting redirection*

“|” (**pipe**)

- Redirecting **output** of one program **as input** of another program

-----

```
$ netstat | less
```

## Utilities: *Pattern matching*

### grep

- “**G**lobally search for a **r**egular **e**xpression and **p**rint matching lines”
- Matches a string/pattern input to files (per line)
  - **-i** to ignore cases
  - **-E** to use (extended) regular expressions (more on this later)
  - **-R** to recursively match

## Utilities: *Pattern matching*

```
$ grep -i "intel" lscpu.out
```

```
Vendor ID:           GenuineIntel
```

```
Model name:          Intel(R) Xeon(R) CPU @ 2.20GHz
```

## Utilities: *File contents*

### **wc**

- “Word count”
- Statistics on bytes, characters, words, lines, etc. of a file

-----

```
$ wc lscpu.out
  24   153 1147 lscpu.out
```

## Utilities: *File contents*

### `diff`

- Difference between two files
- The output shows the differences as --- **which lines** of **file 1** need to be [added, changed, or deleted (**a**, **c**, **d**)] to match **which other lines** of **file 2**



## Utilities: *File contents*

```
$ lscpu -x > lscpux.out
$ diff lscpu.out lscpux.out
5c5
< On-line CPU(s) list: 0-7
---
> On-line CPU(s) mask: ff
15,16c15,16
< CPU MHz:                2200.188
< BogomIPS:               4400.37
```

... see next slide

## Utilities: *File contents*

---

> CPU MHz:

2199.998

> BogoMIPS:

4399.99

23c23

< NUMA node0 CPU(s): 0-7

---

> NUMA node0 CPU(s): ff

# htop or top

```
1 [ 0.8%] Tasks: 25, 38 thr; 1 running
2 [ 1.5%] Load average: 0.88 1.31 1.57
Mem [|||||] 121M/1.91G Uptime: 00:52:01
Swp [ 125M/40.0G]

PID USER      PRI  NI  VIRT   RES   SHR  S  CPU% MEM%   TIME+  Command
1009 root        20   0 1355M 14740 6096 S  0.8  0.7  0:22.68 /src/codevolve/services
1509 root        20   0 25908 2960 2076 R  0.8  0.1  0:00.09 htop
1011 root        20   0 1355M 14740 6096 S  0.0  0.7  0:06.21 /src/codevolve/services
1166 root        20   0 1355M 14740 6096 S  0.0  0.7  0:03.57 /src/codevolve/services
1013 root        20   0 1355M 14740 6096 S  0.0  0.7  0:02.52 /src/codevolve/services
1298 root        20   0 1355M 14740 6096 S  0.0  0.7  0:02.20 /src/codevolve/services
1019 root        20   0 1355M 14740 6096 S  0.0  0.7  0:03.58 /src/codevolve/services
348  rstudio-s    20   0 129M  4476 3644 S  0.0  0.2  0:01.67 /usr/lib/rstudio-server/bin/rsrver
543  rstudio-s    20   0 129M  4476 3644 S  0.0  0.2  0:00.48 /usr/lib/rstudio-server/bin/rsrver
1045 root        20   0 110M  8496 4872 S  0.0  0.4  0:02.51 /usr/bin/x11vnc -noxrecord -xkb -forever -cursor arrow -display :1 -o /var/log/x11vnc.log
1 root        20   0 156M  6280 4016 S  0.0  0.3  0:02.16 /sbin/init
68  root        20   0 94828 11004 10320 S  0.0  0.6  0:00.55 /lib/systemd/systemd-journald
84  root        20   0 33376 2088 1444 S  0.0  0.1  0:00.04 /lib/systemd/systemd-udev
303  systemd-n    20   0 80040 3668 2976 S  0.0  0.2  0:00.07 /lib/systemd/systemd-networkd
306  systemd-r    20   0 70624 3668 3064 S  0.0  0.2  0:00.09 /lib/systemd/systemd-resolved
341  root        20   0 294M  5340 4344 S  0.0  0.3  0:00.09 /usr/lib/accountsservice/accounts-daemon
345  root        20   0 294M  5340 4344 S  0.0  0.3  0:00.00 /usr/lib/accountsservice/accounts-daemon
336  root        20   0 294M  5340 4344 S  0.0  0.3  0:00.13 /usr/lib/accountsservice/accounts-daemon
437  root        20   0 166M 13820 5868 S  0.0  0.7  0:00.00 /usr/bin/python3 /usr/bin/networkd-dispatcher --run-startup-triggers
337  root        20   0 166M 13820 5868 S  0.0  0.7  0:00.15 /usr/bin/python3 /usr/bin/networkd-dispatcher --run-startup-triggers
339  root        20   0 62016 3424 2884 S  0.0  0.2  0:00.13 /lib/systemd/systemd-logind
340  messagebu    20   0 50056 2548 1980 S  0.0  0.1  0:00.13 /usr/bin/dbus-daemon --system --address=systemd: --nofork --nopidfile --systemd-activation --syslog-only
344  daemon      20   0 28328 1264 1056 S  0.0  0.1  0:00.00 /usr/sbin/atd -f
542  rstudio-s    20   0 129M  4476 3644 S  0.0  0.2  0:00.48 /usr/lib/rstudio-server/bin/rsrver
363  root        20   0 16408 1332 1184 S  0.0  0.1  0:00.00 /sbin/agetty -o -p -- \u --noclear --keep-baud console 115200,38400,9600 vt220
385  root        20   0 72292 3544 2816 S  0.0  0.2  0:00.01 /usr/sbin/sshd -D
397  root        20   0 295M  5452 4520 S  0.0  0.3  0:00.00 /usr/lib/policykit-1/polkitd --no-debug
405  root        20   0 295M  5452 4520 S  0.0  0.3  0:00.01 /usr/lib/policykit-1/polkitd --no-debug

F1Help F2Setup F3Search F4Filter F5Free F6SortBy F7Nice F8Nice F9Kill F10Quit
```

## Utilities: *Other commands*

We'll look at a few more shell commands later in the week including:

- `find`
- `sed`
- `make`
- `sort`

... and several others

## Usability: *Navigating your own commands*

- **Up/down arrows** allow you to navigate your prior commands
- **ctrl r** allows **autocompletion** based on command history
- The **history** command is also useful to jog your memory
- Tab completion --- will try to match a file in a specified directory

## Usability: *Navigating your own commands*

- **ctrl c** interrupts and **kills** a process
- **ctrl a** to the **beginning** of the line, **ctrl e** to the **back**
- Stringing commands together using boolean operators (&&) or separators (;)
- **clear** clears the terminal screen

## Usability: *Filesystem navigation tips*

- **pushd** and **popd** --- create a stack of directories
- Setting **variables** to a particular path
  - Backticks ( ``` ) to **evaluate** a command and use its output (often with **pwd**)

-----

```
$ echo `pwd`
```

```
/root/sandbox
```

```
$ s=`pwd` ; echo $s
```

```
/root/sandbox
```

## Examples: *Questions*

1. Does your sandboxed machine support “**avx**” instructions (according to **lscpu**)?
2. How many active sockets (according to **netstat**) are of type “**DGRAM**”?
3. Can you check how many sockets (according to **netstat**) are of type **DGRAM** every 5 seconds? \*\*\*



## Exercises: #1

- `lscpu` has some **output** that we could use as **input** for a **pattern match**

-----

```
$ lscpu | grep "avx"
```

## Examples: #2

- **netstat**'s output shows each active socket (and type) **per line**
- We could use the **output** as **input** for a **pattern match**
- We could then **count** the **number of lines** that matched

-----

```
$ netstat | grep "DGRAM" | wc -l
```

(What's **-l** ? See the manual!)

## Examples: #3

Would be nice to check in **a loop**, maybe use a **script**, and **view** the output **incrementally** ...

We'll start with this tomorrow.