

Crash Course on UNIX and Systems Tools

Day 5 --- More Project Structure and Debugging

Project Structure: *Example makefile*

Here's where we are ... let's improve it

```
CC:=gcc
CFLAGS:=-c
LIBS:=-lm

myprog: lib.o test.o
    $(CC) $^ $(LIBS) -o $@

lib.o: lib.c
    $(CC) $(CFLAGS) $< -o $@

test.o: test.c
    $(CC) $(CFLAGS) $< -o $@

clean:
    rm -f lib.o test.o myprog

.PHONY: clean
```

Project Structure: *Makefiles*

- Let's move our headers and sources into their own directories --- good convention (**`./src`** , **`./include`**)
 - Need to make sure that this is still valid: **`#include "lib.h"`**
- Use '**`-I./myheaders`**' to offer header files to each source

```
gcc -c -I./include lib.c -lm -o lib.o
```

- In the makefile --- simply add the flag at the end of **CFLAGS**

Project Structure: *Makefiles*

```
SRC:=./src
INC:=./include

CC:=gcc
CFLAGS:=-c -I$(INC)
LIBS:=-lm

myprog: lib.o test.o
    $(CC) ^ $(LIBS) -o $@

lib.o: $(SRC)/lib.c
    $(CC) $(CFLAGS) $< -o $@

test.o: $(SRC)/test.c
    $(CC) $(CFLAGS) $< -o $@

clean:
    rm -f lib.o test.o myprog

.PHONY: clean
```

Project Structure: *Makefiles*

- Still not that modularized ... what if we have more `.o` to build?
 - If every `.o` is built in the same way, have one modular rule!
- Use ‘%’ to create a **pattern rule** that targets many files
 - ‘%’ is a **pattern substitution character** --- “%.c” matches all `.c` sources
 - Be careful about targets that can match multiple rules

Project Structure: *Makefiles*

```
SRC:=./src
INC:=./include

CC:=gcc
CFLAGS:=-c -I$(INC)
LIBS:=-lm

myprog: lib.o test.o
    $(CC) $^ $(LIBS) -o $@

%.o: $(SRC)/%.c
    $(CC) $(CFLAGS) $< -o $@

clean:
    rm -f lib.o test.o myprog

.PHONY: clean
```

Project Structure: *Makefiles*

- Still not that modularized ... what about the dependencies?
 - Need to handle growing number of dependencies to build “**myprog**”
- Fetch all the sources at once using the **wildcard** function
 - `SRCS:=$(wildcard $(SRC)/*.c)`
 - Achieve wildcard function in a variable declaration
 - Here, `*.c` matches for all C sources, and the **wildcard** function will fetch these files (perform the expansion) and space-separate each match

Project Structure: *Makefiles*

- We're not done --- notice that the dependencies are `.o`, and we fetched `.c` files (those `.o` files don't exist yet)

```
myprog: lib.o test.o
```

- Create a list of `.o` deps. --- let's use the **patsubst** function
 - `OBJS=$(patsubst $(SRC)/%.c,%.o,$(SRCS))`
`$(patsubst to_match,to_gen,where_to_match_from)`
 - Generates a space-separated string via a pattern and search and replace

Project Structure: *Makefiles*

```
NAME:=myprog
SRC:=./src
INC:=./include

SRCS:=$(wildcard $(SRC)/*.c)
OBJS:=$(subst $(SRC)/%.c,%.o,$(SRCS))

CC:=gcc
CFLAGS:=-c -I$(INC)
LIBS:=-lm
```

```
$(NAME): $(OBJS)
    $(CC) $^ $(LIBS) -o $@

%.o: $(SRC)/%.c
    $(CC) $(CFLAGS) $< -o $@

clean:
    rm -rf $(OBJS) $(NAME)

.PHONY: clean
```

Debugging: *Running myprog*

```
$ ./myprog
```

```
Segmentation fault (core dumped)
```

- **Segmentation fault** --- we're accessing/using memory that we should not have

What now?

Debugging: ***gdb***

- Use a debugger to help you understand the problem --- **gdb**
- First --- compile with **symbols** and **debug information** (**-g**)

```
gcc -c -g -I./include lib.c -o lib.o
```

Debugging: ***gdb***

- Run your program with gdb (with arguments, use the alternative)
 - *Please note in subsequent screenshots --- my paths may not always be the same as yours --- try to locate the right file*

```
$ gdb myprog
```

```
Reading symbols from myprog...done.
```

```
(gdb)
```

Debugging: ***`gdb`***

- `r` to run the program

```
(gdb) r
```

```
Starting program: /root/sandbox/project_structure/myprog
```

```
Program received signal SIGSEGV, Segmentation fault.
```

```
0x00005555555546c2 in recursive_factorial (
```

```
    n=<error reading variable: Cannot access memory at address
```

```
0x7fffffff7feffc>) at src/lib.c:9
```

```
9      {
```

```
(gdb)
```

Debugging: ***gdb***

- **bt** to see a **backtrace** (i.e. the calls that led to this position)

```
(gdb) bt
#0  0x00005555555546c2 in recursive_factorial (
    n=<error reading variable: Cannot access memory at address 0x7ffffff7feffc>) at src/lib.c:9
#1  0x00005555555546d6 in recursive_factorial (n=-54407990) at src/lib.c:13
#2  0x00005555555546d6 in recursive_factorial (n=1287167194) at src/lib.c:13
#3  0x00005555555546d6 in recursive_factorial (n=1090044298) at src/lib.c:13
#4  0x00005555555546d6 in recursive_factorial (n=1785087258) at src/lib.c:13
```

... and keeps going ...

Debugging: ***gdb***

- **list** to see the erroring location in the source code

```
(gdb) list
4
5      /*
6      * Recursive factorial calculation
7      */
8      int recursive_factorial(int n)
9      {
10         /*
11         * Recurse
12         */
13         return recursive_factorial(n * (n - 1));
(gdb)
```

Debugging: ***gdb***

The **conclusion**? We had the following problems:

- An **infinite recursion** caused by **recursive_factorial** (as seen from the stack trace using **bt**)
 - Why a **seg fault**? --- Infinite recursion will “smash the stack” as arguments and return addresses keep piling onto the stack (this is a 213 concept)
- A solution? Add the **base case**!

Debugging: *gdb*

```
5  /*
6  * Recursive factorial calculation
7  */
8  int recursive_factorial(int n)
9  {
10     /*
11     * Handle the base case
12     */
13     if (n == 0) { return 1; }
14
15
16     /*
17     * Recurse
18     */
19     return recursive_factorial(n * (n - 1));
20 }
```

Debugging: ***gdb***

- Let's try it again!

```
$ gdb myprog
Reading symbols from myprog...done.
(gdb) r
Starting program: /root/sandbox/crash/day3/intermediate/myprog

Program received signal SIGSEGV, Segmentation fault.
0x00005555555546c2 in recursive_factorial (
    n=<error reading variable: Cannot access memory at address 0x7ffffff7feffc>) at src/lib.c:9
9      {
```

Same problem again ... let's examine more carefully in ***gdb***

Debugging: ***gdb***

- Let's use a breakpoint --- the program pauses once it reaches a point that we specify from the code. We can examine there
 - Break at **lib.c:19** (according to the stack trace from **bt**)

```
$ gdb myprog
Reading symbols from myprog...done.
(gdb) break lib.c:19
Breakpoint 1 at 0x6d2: file src/lib.c, line 19.
(gdb) r
Starting program: /root/sandbox/crash/day3/intermediate/myprog

Breakpoint 1, recursive_factorial (n=10) at src/lib.c:19
19      return recursive_factorial(n * (n - 1));
(gdb) █
```

Debugging: ***gdb***

- We've hit the breakpoint (the program's execution has reached the line that we've specified)
 - Specifies current function (**`recursive_factorial`**), the specific arguments for this invocation (**`n=10`**), and the line in the source code

```
$ gdb myprog
Reading symbols from myprog...done.
(gdb) break lib.c:19
Breakpoint 1 at 0x6d2: file src/lib.c, line 19.
(gdb) r
Starting program: /root/sandbox/crash/day3/intermediate/myprog

Breakpoint 1, recursive_factorial (n=10) at src/lib.c:19
19      return recursive_factorial(n * (n - 1));
(gdb) █
```

Debugging: ***gdb***

- We can visualize the code again with **list**

```
Breakpoint 1, recursive_factorial (n=10) at src/lib.c:19
19      return recursive_factorial(n * (n - 1));
(gdb) list
14
15
16      /*
17      * Recurse
18      */
19      return recursive_factorial(n * (n - 1));
20  }
21
22
23  /*
```

Debugging: ***gdb***

- We have some info from the breakpoint --- what now?
- **Step** through the code to understand **how it's executing**
 - **step** (or **s**) will execute until the next line of source code and pause
 - Can also step **into a function** using **s** if a certain line has a function call
 - Other instructions like **next** (or **n**) will stay in the current function and step **over** function calls in code

Debugging: ***gdb***

- In this case, we should **step into** since the function is **recursive**
- The following slide shows this process using **s** and **bt** to show where we are as we step into each recursive call

Debugging: ***gdb***

```
Breakpoint 1, recursive_factorial (n=10) at src/lib.c:19
19      return recursive_factorial(n * (n - 1));
(gdb) bt
#0 recursive_factorial (n=10) at src/lib.c:19
#1 0x0000555555554773 in main (argc=1, argv=0x7fffffff308) at src/test.c:8
(gdb) s
recursive_factorial (n=90) at src/lib.c:13
13      if (n == 0) { return 1; }
(gdb) bt
#0 recursive_factorial (n=90) at src/lib.c:13
#1 0x00005555555546e3 in recursive_factorial (n=10) at src/lib.c:19
#2 0x0000555555554773 in main (argc=1, argv=0x7fffffff308) at src/test.c:8
(gdb) s

Breakpoint 1, recursive_factorial (n=90) at src/lib.c:19
19      return recursive_factorial(n * (n - 1));
(gdb) s
recursive_factorial (n=8010) at src/lib.c:13
13      if (n == 0) { return 1; }
(gdb) bt
#0 recursive_factorial (n=8010) at src/lib.c:13
#1 0x00005555555546e3 in recursive_factorial (n=90) at src/lib.c:19
#2 0x00005555555546e3 in recursive_factorial (n=10) at src/lib.c:19
#3 0x0000555555554773 in main (argc=1, argv=0x7fffffff308) at src/test.c:8
(gdb)
```


Debugging: ***gdb***

- By stepping and looking at each recursive call, we see an interesting pattern with the value of **n**:

```
(gdb) bt
#0  recursive_factorial (n=8010) at src/lib.c:13
#1  0x00005555555546e3 in recursive_factorial (n=90) at src/lib.c:19
#2  0x00005555555546e3 in recursive_factorial (n=10) at src/lib.c:19
#3  0x0000555555554773 in main (argc=1, argv=0x7fffffffe308) at src/test.c:8
(gdb)
```

- We're calculating a factorial --- the values of **n** **should be decrementing** --- the issue must be with **which value is passed** into each recursive call

Debugging: ***gdb***

- Revisiting the source code (see previous slides), we see that the “**n ***” in Line 19 is in the completely **wrong place**

```
Breakpoint 1, recursive_factorial (n=10) at src/lib.c:19
19      return recursive_factorial(n * (n - 1));
(gdb) list
14
15
16      /*
17      * Recurse
18      */
19      return recursive_factorial(n * (n - 1));
20  }
21
22
23  /*
```

Debugging: ***gdb***

The fix!

```
/*
 * Recursive factorial calculation
 */
int recursive_factorial(int n)
{
    /*
     * Handle the base case
     */
    if (n == 0) { return 1; }

    /*
     * Recurse
     */
    return n * recursive_factorial(n - 1);
}
```

Debugging: ***`gdb`***

- The program runs, but the output is *suspicious* (see **main** in **test.c**)

```
$ ./myprog
```

```
3628800
```

← 10! according to **recursive_factorial**

```
0
```

← 10! according to **iterative_factorial**

```
0.000000
```

← $\sqrt{10!}$ according to **sqrt_of_factorial** and
iterative_factorial

Debugging: ***`gdb`***

- The output should match --- seems like a problem in `iterative_factorial`

```
$ ./myprog
```

```
3628800
```

← 10! according to `recursive_factorial`

```
0
```

← 10! according to `iterative_factorial`

```
0.000000
```

← $\sqrt{10!}$ according to `sqrt_of_factorial` and `iterative_factorial`

Debugging: ***gdb***

- Let's run under **gdb** and break at **iterative_factorial**

```
$ gdb myprog
Reading symbols from myprog...done.
(gdb) break iterative_factorial
Breakpoint 1 at 0x6ec: file src/lib.c, line 31.
(gdb) r
Starting program: /root/sandbox/crash/day3/intermediate/myprog
3628800

Breakpoint 1, iterative_factorial (n=10) at src/lib.c:31
31          int result = 0;
(gdb)
```

Debugging: ***gdb***

- Since the return value of **iterative_factorial** is 0, we need to see how the factorial is calculated in each step
- We can **print** out a variable from the source code as we step through it

Debugging: ***gdb***

- We can use **n** to step, and seems like **result** is always 0
- Must mean that we initialized **result** poorly --- should be 1

```
Breakpoint 1, iterative_factorial (n=10) at src/lib.c:31
31         int result = 0;
(gdb) n
32         for (int i = 2 ; i <= n ; i++)
(gdb) n
34             result *= i;
(gdb) n
32         for (int i = 2 ; i <= n ; i++)
(gdb) print result
$2 = 0
```


Debugging: ***gdb***

The fix!

```
/*  
 * Iterative factorial calculation  
 */  
int iterative_factorial(int n)  
{  
    /*  
     * Iteratively compute using a loop  
     */  
    int result = 1;  
    for (int i = 2 ; i <= n ; i++)  
    {  
        result *= i;  
    }  
  
    return result;  
}
```

Debugging: ***valgrind***

- Another common tool to debug and profile is **valgrind** --- which specializes in **checking memory usage**
- Very useful to understand:
 - Stack overflows
 - Memory corruption
 - Memory leaks
 - Null pointer dereferences, invalid writes/reads

Debugging: *valgrind*

- Run your executable (with args, flags, etc.) but place “**valgrind**” at the front of the command (**valgrind -v** for verbose checks)

```
$ valgrind ./myprog
==2305== Memcheck, a memory error detector
==2305== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==2305== Using Valgrind-3.13.0 and LibVEX; rerun with -h for copyright info
==2305== Command: ./myprog
==2305==
```

Debugging: *valgrind*

- If we revert changes back to the buggy library and run under **valgrind**, we can see a detailed report of the **stack overflow**:

```
==2361== Stack overflow in thread #1: can't grow stack to 0x1ffe801000
==2361==
==2361== Process terminating with default action of signal 11 (SIGSEGV)
==2361== Access not within mapped region at address 0x1FFE801FF8
==2361== Stack overflow in thread #1: can't grow stack to 0x1ffe801000
==2361==    at 0x1086CD: recursive_factorial (lib.c:19)
==2361== If you believe this happened as a result of a stack
==2361== overflow in your program's main thread (unlikely but
==2361== possible), you can try to increase the size of the
==2361== main thread stack using the --main-stacksize= flag.
==2361== The main thread stack size used in this run was 8388608.
==2361== Stack overflow in thread #1: can't grow stack to 0x1ffe801000
==2361==
==2361== Process terminating with default action of signal 11 (SIGSEGV)
==2361== Access not within mapped region at address 0x1FFE801FF0
==2361== Stack overflow in thread #1: can't grow stack to 0x1ffe801000
==2361==    at 0x4A2A650: _vgnU_freeres (in /usr/lib/valgrind/vgpreload_core-amd64-linux.so)
==2361== If you believe this happened as a result of a stack
==2361== overflow in your program's main thread (unlikely but
==2361== possible), you can try to increase the size of the
==2361== main thread stack using the --main-stacksize= flag.
==2361== The main thread stack size used in this run was 8388608.
```